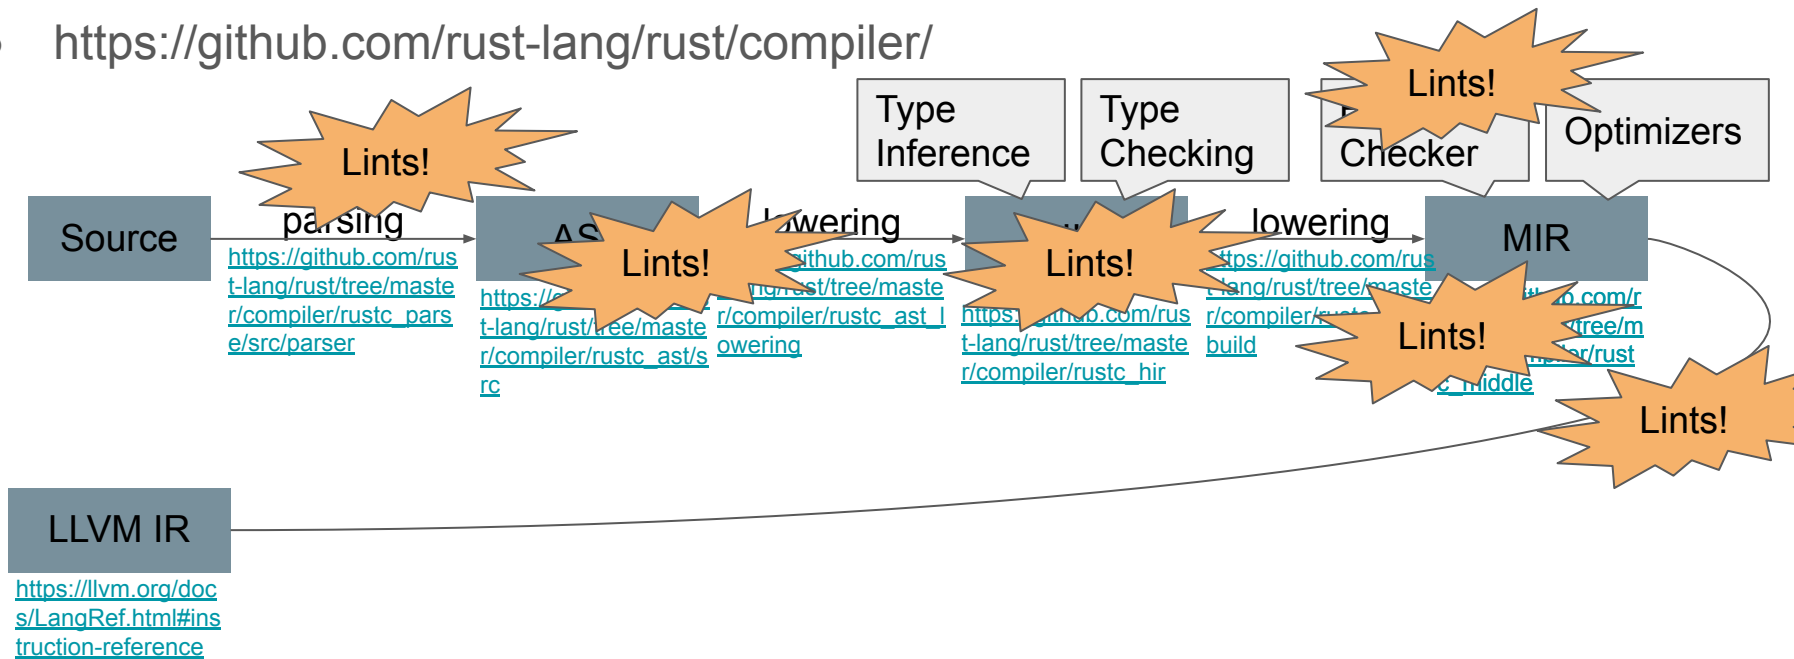


Something interesting about rustc

Polyvios Pratikakis
DevStaff, Oct 13

The Rust Compiler

- <https://github.com/rust-lang/rust/compiler/>



Compiler code base - Interesting stuff #1

- Mem Mgmt tricks (in Rust, so ok)
- Values interned in arenas (Hash-consing)
 - All values allocated in MIR are in the same arena
 - Compare on allocation
 - Pointer is the same \Leftrightarrow Value is the same
 - Arena is lifetime

Compiler code base - Interesting stuff #2

- Query system
 - Be lazy, look like an IDE (but only the good parts)
 - Faster compilation, avoid re-doing work
 - Incremental compilation (no, really)
 - “Do we have optimized MIR for this function?” “Do we have the live variables for this basic block?”
 - “Yes, here” (even for the same file, crate)
 - “No, get me HIR” (this may cascade to another query, and another)
 - Compile lazily
 - Save intermediate state for subsequent compilations
 - If you change only 1 function in a file, why recompile the rest?
- (OK, it's a bit more complex than that.)

Compiler code base - Interesting stuff #3

- Is being parallelized
- Cargo already parallelizes building stuff, but not RustC
- Parallel LLVM code generation
- Other parts of the compiler are being parallelized
 - Not yet there
 - Need people to do the work
 - Good: In Rust, use “async”
 - Good: Borrow checker protects from races!
 - Less good: may need to rewrite lifetimes to partition arenas and then parallelize work

Compiler code base - Interesting stuff #4

- MIR is a control-flow graph (CFG)
- Includes data flow analysis engine
- Pretty easy to write a dataflow analysis
 - (Write your own lints!)
- Pretty easy to write an optimization
 - (Communication with LLVM, well, not so easy)
- Rust has polymorphic code
 - MIR does not
 - No type erasure
 - Instantiation-based monomorphization (like c++)
 - Ew

Historical stuff

- Safe memory management, but fast
 - For many years, science fiction
 - The main problem is `free()`, in C (from ALGOL etc.)
 - Garbage collection (1960)
- 1967 first “Arena”, 1990 in C
- 1988, first idea to get safe “arena”, for objects
- 1994, first theory for safe regions, Standard ML (inference, expensive)
- 2002, Cyclone (C-clone), safe regions (unique, borrowing, lots of annotations)
- Rust: safe region checking (ownership types, borrowed ownership, simpler)
 - Seems to be OK
 - Is there a simpler way? (Except infinite RAM)

Ref

- <https://github.com/rust-lang/rust/tree/master/compiler>
<https://rustc-dev-guide.rust-lang.org/>
- <https://llvm.org/docs/LangRef.html#instruction-reference>