

October 12, 2023

How we made our app scalable

yodeck


Who am I

- Spyros Panagiotopoulos
- Software/DevOps engineer
- Curious about tech
- Problem Solver

Stateful vs Stateless

- Maintains and relies on the state or data between interactions.
 - Uses memory or files to store and retrieve information.
 - Examples include traditional web applications that rely on session cookies to remember user data.
- Does not store session-specific data between interactions.
 - Each request treated as an independent and self-contained transaction.
 - Required data along with request.
 - Highly scalable.

Our CMS instance before



- R5.2xlarge
 - 8 vCPUs
 - 64GB RAM

How we scale?

Monitoring


**If you can't measure it
you can't improve it**

Stand on the shoulders of giants




Consider migration to managed services

- Redis and RabbitMQ with the application server.
 - Makes sense when cost sensitive.
 - Might require scaling along with the application.
- With migration:
 - Pay a little extra
 - Get rid of maintenance




Utilize the `--max-tasks-per-child` flag in Celery

- Celery has issues with memory leaks
 - External libraries make things worse



Utilize the `--max-tasks-per-child` flag in Celery

- Built in flag that kills workers after X tasks
- Find the right tradeoff between memory balance and CPU overhead
 - 400 tasks works well for us



Use a database for Celery Beat

- Celery beat stores tasks metadata in a file by default
 - (Used to be) easily corruptible shelf database



```
1 @periodic_task_that_checks_time(run_every=crontab(hour='0', minute='20', day_of_month=[2, 8, 15]))
2 def retry_partner_unpaid_invoices():
```

Use a database for Celery Beat

- Experimented with django-celery-beat




```
1 @periodic_task_that_checks_time(run_every=crontab(hour='0', minute='20', day_of_month=[2, 8, 15]))
2 def retry_partner_unpaid_invoices():
```

- Ended up with redbeat



```
1 celery -A app beat --detach \  
2     --scheduler=redbeat.RedBeatScheduler
```



RabbitMQ

Redis

celery
Beat

Avoid Writing Data to Files

- State should live outside of the application or its infrastructure
- We stored audit logs in files
 - Rotated each day
 - A task separated logs of each account
- When scaled logs would be missed

Avoid reading data from files

- Don't rely that files generated by your application code would be always there
 - Don't create them in the first place
- Zoho SDK does this by default
 - Persists authentication token in a file
 - Thankfully, base class can be extended

Store sessions in the database

- Web requests are stateless
- To prove identity clients send a piece of information every time
- 2 ways of storing this information
 - Data inside a cookie
 - Data on server, tied with a random string, called Session ID
- Web Servers prefer Session IDs
 - Harder to tamper with
 - Total control of session expiration
 - Store more data than cookie
 - Session data in memory
 - What happens when requests are spread across servers

Store sessions in the database

- Django has an easy way of storing sessions in database
 - Also has an easy way to utilize both database & cache



```
1  INSTALLED_APPS = (  
2      'django.contrib.sessions',  
3  )  
4  MIDDLEWARE = (  
5      'django.contrib.sessions.middleware.SessionMiddleware',  
6  )  
7  SESSION_ENGINE = "django.contrib.sessions.backends.cached_db"
```

Zero downtime rolling updates with uWSGI

uWSGI


- Application server and protocol
 - Serve Python applications
- Single server
 - Pull the code
 - Perform migrations
 - Gracefully reload workers



```
1 touch-chain-reload = /reloadFile
```


uWSGI and Kubernetes

- 2 containers in our pod
- 2 readiness probes
 - Custom nginx endpoint that returns 200
 - GET to our login page




uWSGI and Kubernetes

- A new pod gets created
- When it's ready, it's added in the Kubernetes service
- Starts serving requests



uWSGI and Kubernetes

- What happens when a pod gets terminated?
- Kubernetes sends a SIGTERM signal
 - SIGTERM by default brutally reloads the uWSGI workers



uWSGI and Kubernetes

- uWSGI docs recommend to use the die-on-term flag to respect the SIGTERM signal



- But this brutally kills all the workers

uWSGI and Kubernetes

- SIGTERM signal can be trapped with the following configuration



```
1 hook-master-start = unix_signal:15 gracefully_kill_them_all
```


uWSGI and Kubernetes

- uWSGI exposes a FIFO to interact with it
 - Reload logs
 - Reload the code
 - Gracefully kill workers




```
1  lifecycle:
2    preStop:
3      exec:
4        command:
5          - "/bin/sh"
6          - "-c"
7          - "echo q > /run/uwsgi/uwsgififo"
```

- This should work but we still received 502 responses


uWSGI and Kubernetes

- We use Application Load Balancer from AWS
- When a pod is created and targeted by a service, it is added to the ALB target group
 - We have a health check in ALB as well
- When the pod is ready, ALB sends traffic to it



uWSGI and Kubernetes

- When a pod is terminated, it takes some time to deregister from ALB
- But ALB sends traffic




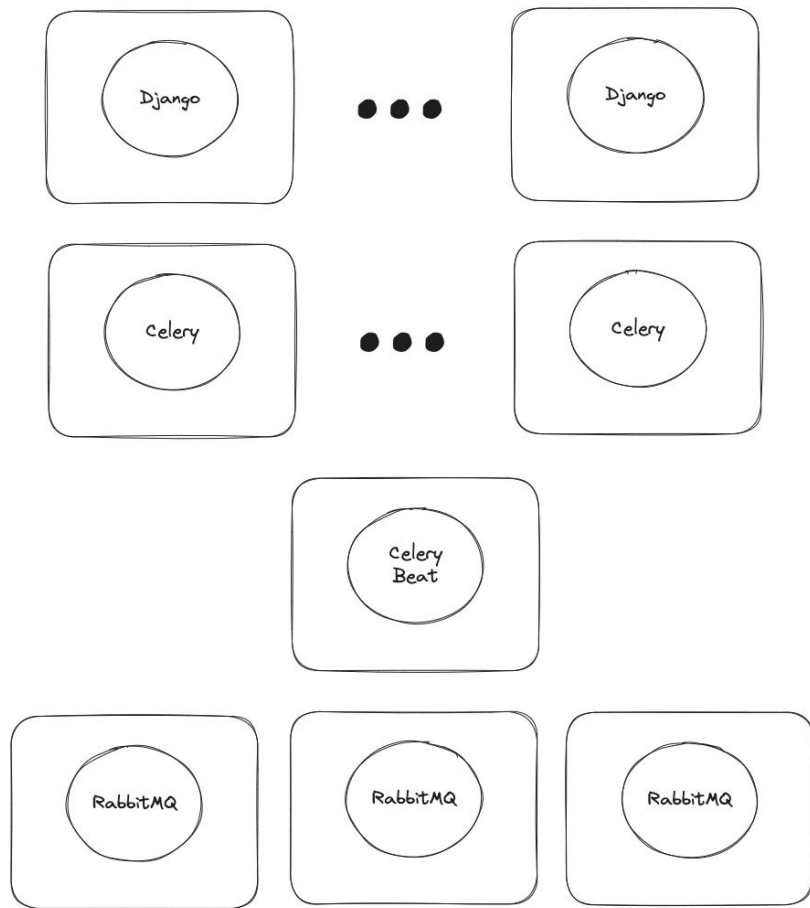
uWSGI and Kubernetes

- We can use **Readiness Gate**
- Injects extra data to the pod
- Kubernetes knows the ALB registration status



```
1 $ kubectl label namespace readiness elbv2.k8s.aws/pod-readiness-gate-inject=enabled
2 namespace/readiness labeled
```

Today



Our CMS now

Questions?

Thank you